

SCHOOL OF COMPUTATION,
INFORMATION AND TECHNOLOGY —
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Master's Thesis in Informatics

**FLOps: Practical Federated Learning via
Automated Orchestration (on the Edge)**

Alexander Malyuk

SCHOOL OF COMPUTATION,
INFORMATION AND TECHNOLOGY —
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Master's Thesis in Informatics

**FLOps: Practical Federated Learning via
Automated Orchestration (on the Edge)**

TODO

Author:	Alexander Malyuk
Supervisor:	Prof. Dr-Ing. Jörg Ott
Advisor:	Dr. Nitinder Mohan, Giovanni Bartolomeo
Submission Date:	15.09.2024

I confirm that this master's thesis is my own work and I have documented all sources and material used.

Munich, 15.09.2024

Alexander Malyuk

Acknowledgments

Abstract

Kurzfassung

Contents

Acknowledgments	iii
Abstract	iv
Kurzfassung	v
1 Introduction	1
1.1 Problem Statement	1
1.2 Motivation	2
1.3 Objectives	3
1.4 Contribution	3
1.5 Thesis Structure	5
2 Background	6
2.1 Federated Learning	6
2.1.1 Basics	6
2.1.2 Architectures	6
2.1.3 Research	6
2.1.4 Industry	6
2.1.5 Frameworks & Libraries	6
2.1.6 Flower	6
2.2 Machine Learning Operations	6
2.2.1 DevOps	6
2.2.2 MLOps	6
2.2.3 MLflow	6
2.3 Orchestration	6
2.3.1 Fundamentals	6
2.3.2 ML Containerization & Orchestration	6
2.3.3 Oakestra	6
2.4 Related Work	6
3 Requirements Analysis	7

4	System Design	8
5	Object Design	9
6	Evaluation	10
6.1	Rationale	10
6.1.1	Chosen Experiments	10
6.2	Experimental Setup	10
6.2.1	Monolith	10
6.2.2	Multi-Cluster	10
6.2.3	Evaluation Procedure	10
6.3	Results	10
6.3.1	Basics	10
6.3.2	Image Builder	10
6.3.3	Different ML Frameworks/Libraries & Datasets	10
6.3.4	Multi-cluster & HFL	10
7	Conclusion	11
7.1	Limitations & Future Work	11
7.1.1	Federated Learning via FLOps	11
7.1.2	Complementary Components & Integrations	11
8	DELME tum example	12
8.1	Section	12
8.1.1	Subsection	12
	Abbreviations	14
	List of Figures	15
	List of Tables	16
	Bibliography	17

1 Introduction

The number of smart devices has been rapidly growing in the last several years. Improvements in connectivity (Cloud Computing & Internet of Things—IoT), connection speeds (5G), and computing power enable this increasing fleet of edge/mobile devices to generate enormous amounts of data (BigData). Combined with the expansion of AI/ML, this data is a driving factor for current successful workflows and future advancements. This complementing union of technologies plays a key role in elevating various domains, from agriculture and healthcare to education and the security sector, to Industry 4.0 and beyond (Industry 5.0). Diverse and complex challenges arise from this swiftly evolving landscape. [1]

1.1 Problem Statement

With great access to data comes great responsibility that can be easily exploited. Many of the aforementioned machines are personal user devices or belong to companies and organizations that handle customer or internal resources. These devices store and handle sensitive private data.

In classical (large-scale) Machine Learning, data gets sent from client devices to a centralized server, which usually resides in the cloud. The collected data is used on the server to train ML models or perform inference serving. This approach provides direct access to this sensitive data and the power to trace back its origin, creating a breach of privacy. [4]

Governments and organizations have established laws and regulations to prohibit potential abuse of sensitive data. Examples include the European Parliament regulation to protect personal data [6] or the California Consumer Privacy Act (CCPA) [3]. These measures aim to support cooperation between organizations and nations while protecting trade secrets. However, some laws and regulations prohibit sharing or moving data to other countries or even off-premises. [4]

Ignoring and no longer using this large amount of data would heavily limit current workflows and further developments for many data-dependent and data-hungry technologies. In 2017, a team of Google researchers introduced Federated Learning (FL) as one possible solution to utilize sensitive data while keeping it private. Instead of collecting the data on a server and training ML models centralized, in FL, the model

training occurs directly on the client devices. Afterward, the individually trained models get sent to the server, which combines the collected models into a single shared one. This so-called global model can then be distributed to the clients again for further training cycles. Therefore, FL enables training a shared model on sensitive data while keeping that data secure on the local client devices. [5]

While many researchers are actively engaged in the field of FL, the majority of them are focused on enhancing existing FL components, strategies, and algorithms or devising better ways of doing FL. However, there is a noticeable scarcity of work that concentrates on the crucial aspects of the initial setup, deployment, and usability of FL. We delve into this issue in greater detail in the dedicated background section (2.1.3).

Because FL is a relatively modern technique, it lacks a sophisticated production-grade ecosystem that includes frameworks and libraries that improve ease of use by automating its setup and execution. As a result, contributing to the field of FL or reproducing findings is a task ranging from non-trivial to improbable due to the lack of documented steps regarding setup, deployment, management, and execution. Instead of using a shared set of tools for bootstrapping to make progress on novel work more efficiently, one needs to set up and manage FL from the ground up. Note that a small set of emerging libraries and frameworks exists for FL. Instead of orchestrating FL on real distributed devices, they focus on executing FL algorithms and processes, often via virtual simulations. Not to mention utilizing more advanced techniques to increase productivity that other domains have already been using for several years, such as modern DevOps practices like continuous deployment. We review existing FL tools in detail in the dedicated section (2.1.5).

1.2 Motivation

Building or contributing to a novel FL framework or library focusing on the previously mentioned challenges areas could soften or entirely alleviate those problems. We are talking about a tool that sees Docker and Kubernetes as role models and strives to be comparable to them but for the discipline of FL. It should specialize in the setup, deployment, component management, and automation, in short, FL orchestration. Allowing researchers, developers, and end-users to set up, perform, reproduce, and experiment with FL in a more accessible way.

The goal of this tool should be to automate and simplify complex tasks, reducing the required level of expertise in various domains, ranging from ML/FL, dependency management, containerization technologies, and automation to orchestration. Such a tool would empower less experienced individuals to participate and contribute to the field of FL. As a result, FL could be adapted and used by more people in more areas.

This tool would streamline and accelerate existing workflows and future progress by utilizing reliable automation to avoid error-prone manual tasks. With its potential to optimize, standardize and unify workflows, our envisioned tool could become a significant part of the emerging FL ecosystem, contributing to the development and progress of the entire field.

1.3 Objectives

This motivation allows us to distill the following key objectives for such a tool.

Improve Accessibility

Making FL more accessible by abstracting away and automating complexities, enables a larger number of individuals to engage with it. Helping FL to expand to more areas will increase its usage and user base, thus raise the general interest and relevance for its field which should aid its development.

Add Automation

Automating tedious, error-prone, and repetitive manual tasks that are necessary to perform FL will free up time and resources for more critical work that leads to further advancements.

Prioritize tangible Applicability

As we discuss in (2.1.3) FL is struggling from a gap between research/virtual-simulation and practical application in real production environments. This tool should focus on being usable in real physical conditions on distributed devices. It should be feasible to incorporate this tool into existing workflows.

Embrace Plasticity

Because FL is such a young field it faces constant change. Naturally, our tool should welcome change in the form of extendability and adaptability. This tool should be flexible and be applicable in a myriad of use cases and scenarios. It should be easy to modify to accommodate evolving needs.

1.4 Contribution

In this thesis we introduce FLOps, a novel open-source foundational work and proof-of-concept that enables real (not simulated) FL to be used, developed and evaluated. FLOps enriches FL with best SOTA practices from the realms of automation, DevOp-

s/MLOps, and orchestration. FLOps is intended to be easy to use for people without previous expert knowledge in FL, MLOps or orchestration. Users can simply provide their ML code in form of a git repository (e.g. in GitHub) as long as this code satisfies some simple structural prerequisites. This repo code gets automatically augmented by FLOps to support FL. FLOps creates a containerized Image with all necessary dependencies to do FL training. These images get automatically build and follow best practices to be as fast and light-weight as possible. These images can be build for multiple different target platforms. Thus enabling FL component images to run e.g. on arm devices like Pis or Nvidia Jetsons. With the help of Oakestra deploys and orchestrates these FL components. FLOps automatically performs FL training. This training process can be observed during runtime via the use of SOTA MLOps tools like MLflow, which offers a sophisticated GUI where users can observe, compare, store, export, share, and organize training runs, metrics, and trained models.

FLOps uses Flower as a FL framework which so far does not support Hierarchical FL. As far as we know, FLOps is the first work that combines Flower and MLflow and allows HFL.

As far as we know, the term FLOps (besides the unit of how powerful something can compute things) has not been used, thus this work is called FLOps but should also open the door into future exiting developments for dedicated ML/Dev-Ops practices specifically for FL.

FLOps also allows to automatically build an inference server/service based on the trained model. This image can then be pulled like any other image by the user and used for arbitrary purposes. If the user requests it FLOps can also directly deploy this trained-model image as an inference service in Oakestra.

Besides the end-user perspective FLOps is intended to be a foundational piece of software that should be easily modifiable and extendable for developers and researchers. We put a lot of effort into writing high quality code, using SOTA libraries and frameworks. We also added many development friendly features into FLOps. E.g. enforced proper styling and typing via formatters and linters, including CI. We created ready-made static images and services that can be used and extended to automate development and evaluation workflows. E.g. we have a Mock-Data-Provider image/service that can act as a data provider to populate the data used for FL training. Or the Inference Service that can be deployed with the base case FLOps project to verify that the trained model works as expected and is able to properly do inference serving. All these images support arm and amd architectures.

We also add base-images with optimal development flags to speed up build and execution times of FLOps so that developers can verify and check their changes more rapidly.

On top of that we also implemented a new CLI tool for Oakestra from scratch that is

used to interact with Oakestra's and FLOps API. Besides that this configurable CLI tool also is capable of visualizing current proceses in a human friendly way in real time as well as trigger evaluation runs and other automated tasks like installing necessary dependencies onto the local machine.

1.5 Thesis Structure

2 Background

2.1 Federated Learning

2.1.1 Basics

2.1.2 Architectures

2.1.3 Research

2.1.4 Industry

2.1.5 Frameworks & Libraries

2.1.6 Flower

2.2 Machine Learning Operations

2.2.1 DevOps

2.2.2 MLOps

2.2.3 MLflow

2.3 Orchestration

2.3.1 Fundamentals

2.3.2 ML Containerization & Orchestration

2.3.3 Oakestra

2.4 Related Work

3 Requirements Analysis

4 System Design

5 Object Design

6 Evaluation

6.1 Rationale

6.1.1 Chosen Experiments

6.2 Experimental Setup

6.2.1 Monolith

6.2.2 Multi-Cluster

6.2.3 Evaluation Procedure

6.3 Results

6.3.1 Basics

6.3.2 Image Builder

6.3.3 Different ML Frameworks/Libraries & Datasets

6.3.4 Multi-cluster & HFL

7 Conclusion

7.1 Limitations & Future Work

7.1.1 Federated Learning via FLOps

7.1.2 Complementary Components & Integrations

8 DELME tum example

8.1 Section

Citation test [2].

Acronyms must be added in `main.tex` and are referenced using macros. The first occurrence is automatically replaced with the long version of the acronym, while all subsequent usages use the abbreviation.

E.g. `\ac{TUM}`, `\ac{TUM}` \Rightarrow Technical University of Munich (TUM), TUM

For more details, see the documentation of the acronym package¹.

8.1.1 Subsection

See Table 8.1, Figure 8.1, Figure 8.2, Figure 8.3.

Table 8.1: An example for a simple table.

A	B	C	D
1	2	1	2
2	3	2	3




Figure 8.1: An example for a simple drawing.

¹<https://ctan.org/pkg/acronym>




Figure 8.2: An example for a simple plot.

```
SELECT * FROM tbl WHERE tbl.str = "str"
```

Figure 8.3: An example for a source code listing.

Abbreviations

TUM Technical University of Munich

List of Figures

8.1	Example drawing	12
8.2	Example plot	13
8.3	Example listing	13

List of Tables

8.1	Example table	12
-----	-------------------------	----

Bibliography

- [1] A. Bourechak, O. Zedadra, M. N. Kouahla, A. Guerrieri, H. Seridi, and G. Fortino. "At the Confluence of Artificial Intelligence and Edge Computing in IoT-Based Applications: A Review and New Perspectives." In: *Sensors* 23.3 (2023). issn: 1424-8220. DOI: 10.3390/s23031639.
- [2] L. Lamport. *LaTeX : A Documentation Preparation System User's Guide and Reference Manual*. Addison-Wesley Professional, 1994.
- [3] C. Legislature. *California Consumer Privacy Act (CCPA)*. Online; accessed August 11, 2024. 2018.
- [4] H. Ludwig and N. Baracaldo, eds. *Federated Learning - A Comprehensive Overview of Methods and Applications*. Springer, 2022. ISBN: 978-3-030-96896-0. DOI: 10.1007/978-3-030-96896-0.
- [5] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y. Arcas. "Communication-Efficient Learning of Deep Networks from Decentralized Data." In: *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics*. Ed. by A. Singh and J. Zhu. Vol. 54. Proceedings of Machine Learning Research. PMLR, 20–22 Apr 2017, pp. 1273–1282.
- [6] T. E. Parliament and Council. *Regulation (EU) 2016/679 of the European Parliament and of the Council of 27 April 2016 on the protection of natural persons with regard to the processing of personal data and on the free movement of such data, and repealing Directive 95/46/EC (General Data Protection Regulation)*. Online; accessed August 11, 2024. 2016.